

OKE lab completion sheet

Record evidence at the existing checkpoints - no extra exercise.

Keep credentials, kubeconfig contents, and tokens out of this sheet.

Name: _____ Date: _____ Tested repository revision: _____

Checkpoints

- ☐ Preflight passed for my assigned context; two workers are Ready with numeric resource metrics.
- ☐ My public app returned my customized message; two app pods showed 2/2 Ready.
- ☐ I identified web and istio-proxy and found hello-oke-traffic -> hello-oke in Kiali.
- ☐ Manual scaling changed app replicas 2 -> 4 -> 2; worker count and Service IP stayed unchanged.
- ☐ HPA utilization became numeric; I observed scale-out above two and scale-in back to two.
- ☐ I explicitly reset traffic.loadEnabled=false and confirmed two Ready app pods.

My customized response message:

Observations

Use Grafana's Last 30 minutes range and p95 latency for each row. Record HPA replicas with `kubectl get hpa` and Ready app pods with `kubectl get pods`; the Grafana proxy count is not a readiness check. Write no data if a panel is empty, rather than treating it as zero.

Phase	HPA replicas / Ready app pods	Requests/s	Success %	p95 latency (ms)	App proxies up
Baseline (/)	Not enabled / ____	____	____	____	____
During load (/work)	____ / ____	____	____	____	____
After scale-in (/)	____ / ____	____	____	____	____

Peak HPA replica count observed: ____ Time load started: ____ Time scale-in finished: ____

Six replicas is a limit, not a required peak. The two endpoints do different work; this comparison does not isolate autoscaling's effect on latency.

OKE lab completion sheet

Five-minute debrief - explain what you observed.

Name: _____ Date: _____

1. Which Kubernetes object keeps the application reachable as pods change?

2. Which component supplies this HPA's CPU metrics, and which supplies dashboard data?

3. What changed during scaling: app replicas, worker nodes, Service IP?

4. What proves scale-in finished? Why can Grafana's proxy count lag?

5. Which probe removes an unready pod from normal Service traffic, and which can restart a container? Does Kiali's Degraded label identify a failed probe?

Prediction from step 5: a 200m CPU request with a 60% HPA target corresponds to ____ CPU usage.

Optional pod recovery: skipped / completed / blocked. Replacement pod name, if observed: _____

Result: completed / needs instructor help / demonstration only.

If blocked, record the step, symptom, and last observed state:

Leave releases installed for Luna's session-end cleanup. Stop local watches and port-forwards after recording your results; doing so does not stop in-cluster traffic. Give the instructor your blocked checkpoint before leaving.